



# SQL

## Lecture #4

# Setting the Context Using Project Stage 2





*\$sqlite3*

*SQLite version 3.37.2 2022-01-06*

*13:25:41*

*Enter ".help" for instructions*

*Enter SQL statements terminated*

*with a ";"*

*sqlite>*

```
CREATE TABLE regions (
region_id INTEGER PRIMARY KEY AUTOINCREMENT NOT
NULL,
region_name text NOT NULL
);
```

```
CREATE TABLE countries (
country_id text NOT NULL,
country_name text NOT NULL,
region_id INTEGER NOT NULL,
PRIMARY KEY (country_id ASC),
FOREIGN KEY (region_id) REFERENCES regions
(region_id) ON DELETE CASCADE ON UPDATE CASCADE
);
```

```
CREATE TABLE locations (
location_id INTEGER PRIMARY KEY AUTOINCREMENT NOT
NULL,
street_address text,
postal_code text,
city text NOT NULL,
state_province text,
country_id INTEGER NOT NULL,
FOREIGN KEY (country_id) REFERENCES countries
(country_id) ON DELETE CASCADE ON UPDATE CASCADE
);
```

Now load the 7 csv files into the above empty tables. To do so, execute the following steps:

1. Run the following command in SQLite

```
sqlite> .mode csv
```

2. Now run the following command 7 times to load data into the 7 tables. Here 'filename' is the name of the csv file.

```
sqlite> .import '/tail -n +2 {filepath}/{filename}.csv'  
{filename}
```

For example, suppose the 7 csv files are stored in directory /users/user1/cs564 on the machine that you are using, then *sqlite> .import '/tail -n +2 /Users/user1/CS564\_stage2/countries.csv' countries* will load data from countries.csv into table 'countries' in SQLite.

<b>Table</b>	<b>Rows</b>
employees	40
dependents	30
departments	11
jobs	19
locations	7
countries	25
regions	4

# SQL Introduction

Standard language for querying and manipulating data

## Structured Query Language

Many standards out there: SQL92, SQL2, SQL3, SQL99

Vendors support various subsets of these, but all of what we'll be talking about.

# The Main Thing to Remember

```
SELECT    S  
FROM      R1,...,Rn  
WHERE     C1  
GROUP BY a1,...,ak  
HAVING   C2
```

- Later: the devil is in the details

# Select-From-Where Statements

**SELECT**    desired attributes

**FROM**      one or more tables

**WHERE**     condition about tuples of  
                  the tables

# Single-Table Queries

# Our Running Example

- Most of our SQL queries will be based on the following database schema.
  - Underline indicates key attributes.

Beers(name, manf)

Bars(name, addr, license)

Drinkers(name, addr, phone)

Likes(drinker, beer)

Sells(bar, beer, price)

Frequents(drinker, bar)

# Example

- Using Beers(name, manf), what beers are made by Anheuser-Busch?

```
SELECT name  
FROM Beers  
WHERE manf = 'Anheuser-Busch' ;
```

Beers(name	manf)
Bud	Anheuser-Busch
Bud Lite	Anheuser-Busch
Michelob	Anheuser-Busch
Spotted Cow	New Glarus

# Result of Query

name	Beers(name)	manf)
‘Bud’	Bud	Anheuser-Busch
‘Bud Lite’	Bud Lite	Anheuser-Busch
‘Michelob’	Michelob	Anheuser-Busch
	Spotted Cow	New Glarus

The answer is a relation with a single attribute, name, and tuples with the name of each beer by Anheuser-Busch, such as Bud.

# Meaning of Single-Relation Query

- Begin with the relation in the FROM clause.
- Apply the selection indicated by the WHERE clause.
- Apply the extended projection indicated by the SELECT clause.

# Operational Semantics

- To implement this algorithm think of a *tuple variable* ranging over each tuple of the relation mentioned in FROM.
- Check if the “current” tuple satisfies the WHERE clause.
- If so, compute the attributes or expressions of the SELECT clause using the components of this tuple.

## \* In SELECT clauses

- When there is one relation in the FROM clause,
  - \* in the SELECT clause stands for “all attributes of this relation.”
- Example using Beers(name, manf):

```
SELECT *
```

```
FROM Beers
```

```
WHERE manf = 'Anheuser-Busch' ;
```

# Result of Query:

name	manf
‘Bud’	‘Anheuser-Busch’
‘Bud Lite’	‘Anheuser-Busch’
‘Michelob’	‘Anheuser-Busch’

Now, the result has each of the attributes of Beers.

# Another Example

Company(sticker, name, country, stockPrice)

Find all US companies whose stock is > 50:

```
SELECT *
FROM Company
WHERE country="USA" AND stockPrice > 50
```

Output schema: R(sticker, name, country, stockPrice)

# Renaming Attributes

- If you want the result to have different attribute names, use “AS <new name>” to rename an attribute.
- Example based on Beers(name, manf):

```
SELECT name AS beer, manf  
FROM Beers  
WHERE manf = 'Anheuser-Busch'
```

# Result of Query:

beer	manf
‘Bud’	‘Anheuser-Busch’
‘Bud Lite’	‘Anheuser-Busch’
‘Michelob’	‘Anheuser-Busch’

# Expressions in SELECT Clauses

- Any expression that makes sense can appear as an element of a SELECT clause.
- Example: from Sells(bar, beer, price):

```
SELECT bar, beer,  
      price * 120 AS priceInYen  
FROM Sells;
```

# Result of Query

bar	beer	priceInYen
Joe's	Bud	300
Sue's	Miller	360
...	...	...

# Another Example: Constant Expressions

- From Likes(drinker, beer):

```
SELECT drinker,  
       'likes Bud' AS whoLikesBud  
FROM Likes  
WHERE beer = 'Bud' ;
```

# Result of Query

drinker	whoLikesBud
‘Sally’	‘likes Bud’
‘Fred’	‘likes Bud’
...	...

# Complex Conditions in WHERE Clause

- From Sells(bar, beer, price), find the price Joe's Bar charges for Bud:

```
SELECT price  
FROM Sells  
WHERE bar = 'Joe''s Bar' AND  
      beer = 'Bud';
```

# Selections

## What you can use in WHERE:

attribute names of the relation(s) used in the FROM.

comparison operators:  $=$ ,  $\neq$ ,  $<$ ,  $>$ ,  $\leq$ ,  $\geq$

apply arithmetic operations:  $\text{stockprice}^*2$

operations on strings (e.g., “ $\|$ ” for concatenation).

Lexicographic order on strings.

Pattern matching:  $s \text{ LIKE } p$

Special stuff for comparing dates and times.

# Important Points

- Two single quotes inside a string represent the single-quote (apostrophe).
- Conditions in the WHERE clause can use AND, OR, NOT, and parentheses in the usual way boolean conditions are built.
- SQL is *case-insensitive*. In general, upper and lower case characters are the same, except inside quoted strings.

# Patterns

- WHERE clauses can have conditions in which a string is compared with a pattern, to see if it matches.
- General form: <Attribute> LIKE <pattern>  
or <Attribute> NOT LIKE <pattern>
- Pattern is a quoted string with % = “any string”; \_ = “any character.”

# Example

- From Drinkers(name, addr, phone) find the drinkers with exchange 555:

```
SELECT name  
FROM Drinkers  
WHERE phone LIKE '%555-__-__-' ;
```

# The **LIKE** operator

- $s \text{ LIKE } p$ : pattern matching on strings
- $p$  may contain two special symbols:
  - $\%$  = any sequence of characters
  - $_$  = any single character

Company(sticker, name, address, country, stockPrice)

Find all US companies whose address contains “Mountain”:

```
SELECT *
FROM Company
WHERE country="USA" AND
address LIKE "%Mountain%"
```

# Multi-Table Queries

# Multirelation Queries

- Interesting queries often combine data from more than one relation.
- We can address several relations in one query by listing them all in the FROM clause.
- Distinguish attributes of the same name by “<relation>.<attribute>”

# Example

- Using relations Likes(drinker, beer) and Frequent(drinker, bar), find the beers liked by at least one person who frequents Joe's Bar.

```
SELECT beer  
FROM Likes, Frequent  
WHERE bar = 'Joe''s Bar' AND  
Frequent.drinker = Likes.drinker;
```

Likes(drinker, beer)

a	x
b	y
c	z

Find beers liked by at least one person  
who frequents bar t.

Select beer

From Likes, Frequent

Where bar = 't' and

Likes.drinker = Frequent.drinker

Frequent(drinker, bar)

a	t
b	u
c	t

Likes(drinker, beer)

a	x
b	y
c	z

Select beer  
From Likes, Frequents  
Where bar = 't' and  
Likes.drinker = Frequents.drinker

Frequents(drinker, bar)

a	t
b	u
c	t

Solution:

- 1) Enumerate all combinations of tuples from Likes and Frequents
- 2) Keep only combinations that satisfy the condition in Where clause
- 3) Return beer from those combinations

Likes(drinker, beer)

a	x
b	y
c	z

Frequents(drinker, bar)

a	t
b	u
c	t

Select beer  
From Likes, Frequents  
Where bar = 't' and  
**Likes.drinker = Frequents.drinker**

Solution:

- 1) Enumerate all combinations of tuples from Likes and Frequents
- 2) Keep only combinations that satisfy the condition in Where clause
- 3) Return beer from those combinations

(a,x) (a,t)

(a,x) (b,u)

(a,x) (c,t)

(b,y) (a,t)

(b,y) (b,u)

(b,y) (c,t)

(c,z) (a,t)

(c,z) (b,u)

(c,z) (c,t)

(a,x) (a,t)

(a,x) (b,u)

(a,x) (c,t)

(b,y) (a,t)

(b,y) (b,u)

(b,y) (c,t)

(c,z) (a,t)

(c,z) (b,u)

(c,z) (c,t)

Output(beer)

x

z

Likes(drinker, beer)

a	x
b	y
c	z

Frequents(drinker, bar)

a	t
b	u
c	t

Select beer  
From Likes, Frequents  
Where **bar = 't'** and  
**Likes.drinker = Frequents.drinker**

Solution:

- 1) Enumerate all combinations of tuples from Likes and Frequents
- 2) Keep only combinations that satisfy the condition in Where clause
- 3) Return beer from those combinations

**How to do this fast? (Will cover this later in the course)**

1. Have an index on bar for table Frequents
2. Use the index to quickly find tuples in Frequents with bar = t
3. Find drinkers in these tuples (a and c in this case)
4. Have an index on drinker for table Likes
5. Use the index to quickly find tuples in Likes with drinker a, c
6. Find bars in these tuples and return (x, z)

Many different ways to do this, some may be much faster than others

# Another Example

Product (pname, price, category, maker)

Purchase (buyer, seller, store, product)

Company (cname, stockPrice, country)

Person(pname, phoneNumber, city)

Find names of people living in Champaign that bought gizmo products, and the names of the stores they bought from

```
SELECT    pname, store
FROM      Person, Purchase
WHERE     pname=buyer AND city="Champaign"
          AND product="gizmo"
```

# Disambiguating Attributes

Find names of people buying telephony products:

Product (name, price, category, maker)

Purchase (buyer, seller, store, product)

Person(name, phoneNumber, city)

```
SELECT Person.name  
FROM Person, Purchase, Product  
WHERE Person.name=Purchase.buyer  
      AND product=Product.name  
      AND Product.category="telephony"
```

# Disambiguating Attributes

Find names of people buying telephony products:

Product (name, price, category, maker)

Purchase (buyer, seller, store, product)

Person(name, phoneNumber, city)

```
SELECT Person.name  
FROM Person x, Purchase y, Product z  
WHERE x.name=y.buyer  
      AND y.product=z.name  
      AND z.category="telephony"
```


# Formal Semantics

- Almost the same as for single-relation queries:
  1. Start with the product of all the relations in the FROM clause.
  2. Apply the selection condition from the WHERE clause.
  3. Project onto the list of attributes and expressions in the SELECT clause.

# Operational Semantics

- Imagine one tuple-variable for each relation in the FROM clause.
  - These tuple-variables visit each combination of tuples, one from each relation.
- If the tuple-variables are pointing to tuples that satisfy the WHERE clause, send these tuples to the SELECT clause.

# Example



# Explicit Tuple-Variables

- Sometimes, a query needs to use two copies of the same relation.
- Distinguish copies by following the relation name by the name of a tuple-variable, in the FROM clause.
- It's always an option to rename relations this way, even when not essential.

# Example

- From Beers(name, manf), find all pairs of beers by the same manufacturer.
  - Do not produce pairs like (Bud, Bud).
  - Produce pairs in alphabetic order, e.g. (Bud, Miller), not (Miller, Bud).

```
SELECT b1.name, b2.name  
FROM Beers b1, Beers b2  
WHERE b1.manf = b2.manf AND  
      b1.name < b2.name;
```

# Tuple Variables

Find pairs of companies making products in the same category

```
SELECT product1.maker, product2.maker  
FROM Product AS product1, Product AS product2  
WHERE product1.category=product2.category  
AND product1.maker <> product2.maker
```

Product ( name, price, category, maker)

# Tuple Variables

Tuple variables introduced automatically by the system:

Product ( name, price, category, maker)

Becomes:

```
SELECT name  
FROM Product  
WHERE price > 100
```

```
SELECT Product.name  
FROM Product AS Product  
WHERE Product.price > 100
```

Doesn't work when Product occurs more than once:

In that case the user needs to define variables explicitly.

# Meaning (Semantics) of SQL Queries

**SELECT** a<sub>1</sub>, a<sub>2</sub>, ..., a<sub>k</sub>  
**FROM** R<sub>1</sub> AS x<sub>1</sub>, R<sub>2</sub> AS x<sub>2</sub>, ..., R<sub>n</sub> AS x<sub>n</sub>  
**WHERE** Conditions

1. Nested loops:

```
Answer = {}  
for x1 in R1 do  
    for x2 in R2 do  
        ....  
        for xn in Rn do  
            if Conditions  
            then Answer = Answer U {(a1,...,ak)}  
return Answer
```

# Meaning (Semantics) of SQL Queries

```
SELECT a1, a2, ..., ak  
FROM   R1 AS x1, R2 AS x2, ..., Rn AS xn  
WHERE  Conditions
```

## 2. Parallel assignment

```
Answer = {}  
for all assignments x1 in R1, ..., xn in Rn do  
    if Conditions then Answer = Answer U {(a1,...,ak)}  
return Answer
```

Doesn't impose any order !

# Tip for Writing SQL Queries

Product (pname, price, category, maker)

Purchase (buyer, seller, store, product)

Company (cname, stockPrice, country)

Person(pname, phoneNumber, city)

Find names of people living in Champaign that bought gizmo products, and the names of the stores they bought from

```
SELECT    pname, store  
FROM      Person, Purchase  
WHERE    pname=buyer AND city="Champaign"  
          AND product="gizmo"
```

# Exercises

Product ( pname, price, category, maker)

Purchase (buyer, seller, store, product)

Company (cname, stock price, country)

Person( per-name, phone number, city)

**Ex #1:** Find people who bought telephony products.

**Ex #2:** Find names of people who bought American products

**Ex #3:** Find names of people who bought American products and did not buy French products

**Ex #4:** Find names of people who bought American products and they live in Champaign.

**Ex #5:** Find people who bought stuff from Joe or bought products from a company whose stock prices is more than \$50.

# Aggregation

# Aggregations

- SUM, AVG, COUNT, MIN, and MAX can be applied to a column in a SELECT clause to produce that aggregation on the column.
- Also, COUNT(\*) counts the number of tuples.

# Example: Aggregation

- From Sells(bar, beer, price), find the average price of beer x:

```
SELECT AVG(price)
```

```
FROM Sells
```

```
WHERE beer = 'x' ;
```

Sells(bar, beer, price)

t, x, 5

t, y, 4  5  
u, x, 7  7  6

u, y, 8

# Group-by and Having

# Group-By and Having

```
SELECT    S  
FROM      R1,...,Rn  
WHERE     C1  
GROUP BY a1,...,ak  
HAVING   C2
```

1. Enumerate all combinations of tuples from R<sub>1</sub>, ..., R<sub>n</sub>
2. Keep only combinations satisfying condition C<sub>1</sub>
3. Group them by a<sub>1</sub>, ..., a<sub>k</sub>
4. Keep only groups satisfying condition C<sub>2</sub>
5. Pull out from each group the values requested in S  
+ if any aggregation, then apply within the group

Likes(drinker, beer)

a	x
b	y
c	z

For each beer find the minimum price that it is being sold in a bar

Frequents(drinker, bar)

a	t
b	u
c	t

Select beer, min(price)  
From Sells  
Groupby beer

(t, x, 5)  
(u, x, 7)

Sells(bar, beer, price)

t,	x,	5
t,	y,	4
u,	x,	7
u,	y,	8

(t, y, 4)  
(u, y, 8)

Output table: (x, 5)  
(y, 4)

Likes(drinker, beer)

a	x
b	y
c	z

For each bar find the average price at which it sells beers

Frequents(drinker, bar)

a	t
b	u
c	t

Select bar, avg(price)  
From Sells  
Groupby bar

(t, x, 5)  
(t, y, 4)

Sells(bar, beer, price)

t,	x,	5
t,	y,	4
u,	x,	7
u,	y,	8

(u, x, 7)  
(u, y, 8)

Output table: (t, 4.5)  
(u, 7.5)

Likes(drinker, beer)

a	x
b	y
c	z

Find all bars where the avg price of beer exceeds \$5 and then list the most expensive price for each such bar.

Frequents(drinker, bar)

a	t
b	u
c	t

Select bar, max(price)  
From Sells  
Groupby bar  
Having avg(price) > 5

Sells(bar, beer, price)

t,	x,	5
t,	y,	4
u,	x,	7
u,	y,	8

(t, x, 5)  
(t, y, 4)

(u, x, 7)  
(u, y, 8)

Output table: (u, 8)

Likes(drinker, beer)

a	x
b	y
c	z

Find all bars that sell only beer x or y, and where the avg price of beer exceeds \$5, and then list the most expensive price for each such bar.

Frequents(drinker, bar)

a	t
b	u
c	t

Select bar, max(price)

From Sells

Where beer = x OR beer = y

Groupby bar

Having avg(price) > 5

Sells(bar, beer, price)

t,	x,	5
t,	y,	4
u,	x,	7
u,	y,	8
v,	z,	100

(t, x, 5)

(t, y, 4)

(u, x, 7)

(u, y, 8)

Output table: (u, 8)

Likes(drinker, beer)

a	x
b	y
c	z

For each drinker find the average price of beer x at the bars they frequent

Frequents(drinker, bar)

a	t
b	u
<b>b</b>	<b>t</b>

Select drinker, avg(price)  
From Frequents F, Sells S  
Where F.bar = S.bar and S.beer = x  
Groupby drinker

Sells(bar, beer, price)

t,	x,	5
t,	y,	4
u,	x,	7
u,	y,	8
v,	z,	100

(a,t) (t,x,5)  
(b,u) (u,x,7)  
(b,t) (t,x,5)

Output table: (a, 5)  
(b, 6)

# “Common Sense” Requirements

```
SELECT    S  
FROM      R1,...,Rn  
WHERE     C1  
GROUP BY a1,...,ak  
HAVING   C2
```

S = may contain attributes a<sub>1</sub>,...,a<sub>k</sub> and/or any aggregates but NO OTHER ATTRIBUTES

C1 = is any condition on the attributes in R<sub>1</sub>,...,R<sub>n</sub>

C2 = is any condition on aggregate expressions

# General form of Grouping and Aggregation

```
SELECT S  
FROM R1,...,Rn  
WHERE C1  
GROUP BY a1,...,ak  
HAVING C2
```

Evaluation steps:

1. Compute the FROM-WHERE part, obtain a table with all attributes in R<sub>1</sub>,...,R<sub>n</sub>
2. Group by the attributes a<sub>1</sub>,...,a<sub>k</sub>
3. Compute the aggregates in C<sub>2</sub> and keep only groups satisfying C<sub>2</sub>
4. Compute aggregates in S and return the result