


File Organization and Indexing

Lecture 7

Motivation


- ❖ Frequent operations
 - scan (go over all tuples)
 - sort, equality search, range search
 - insert tuple, delete tuple (modify tuple?)
- ❖ Need to speed up these operations
 - using certain file organizations and indexes
- ❖ File organizations
 - heap
 - sorted
 - hash
 - B+

Heap File Implemented as a List



- ❖ The header page id and Heap file name must be stored someplace.
- ❖ Each page contains 2 `pointers' plus data.

Heap File Using a Page Directory



- ❖ The entry for a page can include the number of free bytes on the page.
- ❖ The directory is a collection of pages; linked list implementation is just one alternative.



Hash Files

- ❖ n buckets
- ❖ Each bucket = linked list of pages
- ❖ Hash each tuple into a bucket


Sorted Files

- ❖ Sort on a single attribute
- ❖ Or on a combination of attributes
- ❖ These are called “search keys” or just “keys”
 - not to be confused with keys of an entity set or table

Sorted Files




Sorted Files



B+ Tree Files

- ❖ Sorted file with a lot of pointers on top to direct search




Discussion

- ❖ Frequent operations
 - scan
 - sort
 - equality / range search
 - insert / delete tuples

Motivation for Indexes

- ❖ Frequent operations
 - search by person name
 - search by age, sal, or (age, sal)




Types of Indexes

- ❖ Clustered vs unclustered indexes
- ❖ Primary vs secondary indexes

B+ Tree Index


B+ Tree: Most Widely Used Index

- ❖ Insert/delete at $\log_F N$ cost; keep tree *height-balanced*. (F = fanout, N = # leaf pages)
- ❖ Minimum 50% occupancy (except for root). Each node contains $d \leq m \leq 2d$ entries. The parameter d is called the *order* of the tree.
- ❖ Supports equality and range-searches efficiently.




Example B+ Tree

- ❖ $d = 2$, each entry is a number



Example B+ Tree

- ❖ Search begins at root, and key comparisons direct it to a leaf
- ❖ Search for 5*, 15*, all data entries $\geq 24^*$...



□ Based on the search for 15*, we know it is not in the tree!

B+ Trees in Practice

- ❖ Typical order: 100. Typical fill-factor: 67%.
 - average fanout = 133
- ❖ Typical capacities:
 - Height 4: $133^4 = 312,900,700$ records
 - Height 3: $133^3 = 2,352,637$ records
- ❖ Can often hold top levels in buffer pool:
 - Level 1 = 1 page = 8 Kbytes
 - Level 2 = 133 pages = 1 Mbyte
 - Level 3 = 17,689 pages = 133 MBytes

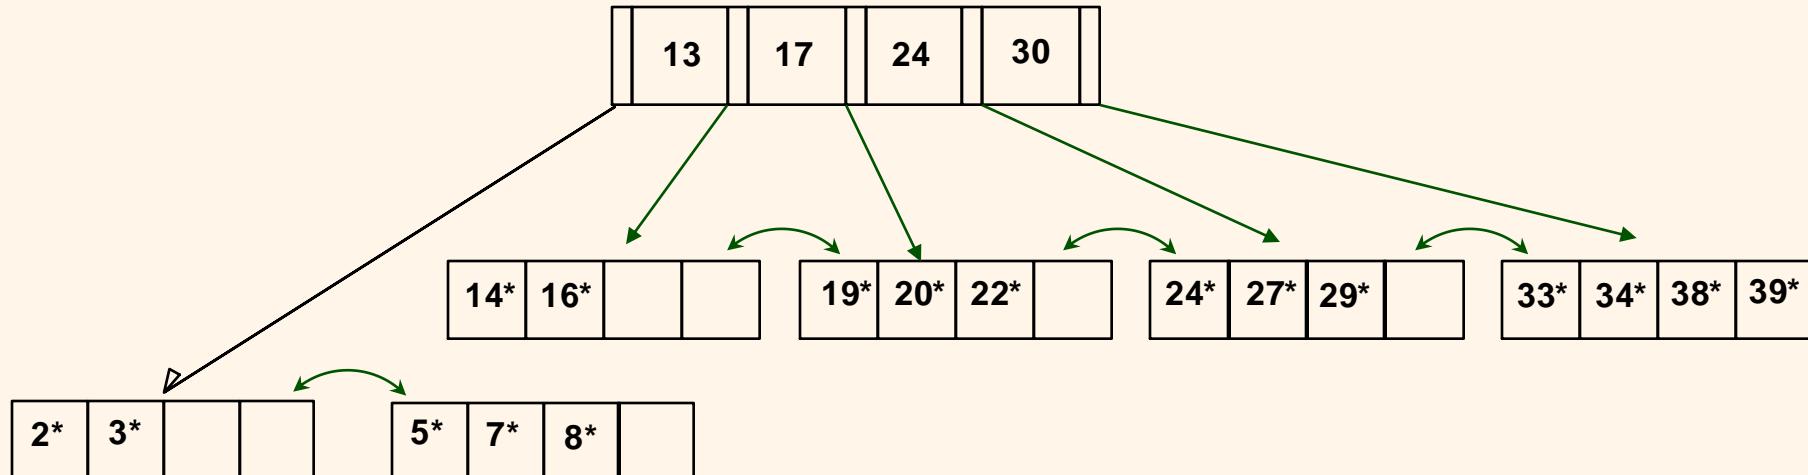

Inserting a Data Entry into a B+ Tree


- ❖ Find correct leaf L .
- ❖ Put data entry onto L .
 - If L has enough space, *done!*
 - Else, must split L (*into L and a new node $L2$*)
 - Redistribute entries evenly, copy up middle key.
 - Insert index entry pointing to $L2$ into parent of L .
- ❖ This can happen recursively
 - To split **index node**, redistribute entries evenly, but push up middle key. (Contrast with leaf splits.)
- ❖ Splits “grow” tree; root split increases height.
 - Tree growth: gets wider or one level taller at top.


Important Tip to Get This Done Right

- ❖ on the exam
- ❖ focus on creating the pages, get them right
- ❖ don't worry about the pointers
- ❖ once you have the pages right, you can easily create the pointers

*Example: Insert 8**







Need to split the root page into two pages
and push up the middle key into another page

17

5 13

24 30

14*	16*			
-----	-----	--	--	--

19*	20*	22*	
-----	-----	-----	--


24*	27*	29*	
-----	-----	-----	--

33*	34*	38*	39*
-----	-----	-----	-----

2*	3*		
----	----	--	--

5*	7*	8*	
----	----	----	--


*Example B+ Tree After Inserting 8**



- ❖ Notice that root was split, leading to increase in height.

Inserting 8 into Example B+ Tree*


- ❖ Observe how minimum occupancy is guaranteed in both leaf and index pg splits.
- ❖ Note difference between *copy-up* and *push-up*; be sure you understand the reasons for this.



Deleting a Data Entry from a B+ Tree

- ❖ Start at root, find leaf L where entry belongs.
- ❖ Remove the entry.
 - If L is at least half-full, *done!*
 - If L has only **d-1** entries,
 - Try to **re-distribute**, borrowing from *sibling* (*adjacent node with same parent as L*).
 - If re-distribution fails, **merge** L and sibling.
- ❖ If merge occurred, must delete entry (pointing to L or sibling) from parent of L .
- ❖ Merge could propagate to root, decreasing height.

Example Tree After (Inserting 8*, Then) Deleting 19* and 20* ...



- ❖ Deleting 19* is easy.
- ❖ Deleting 20* is done with re-distribution.
Notice how middle key is *copied up*.

... And Then Deleting 24*

- ❖ Must merge.
- ❖ Observe *'toss'* of index entry (on right), and *'pull down'* of index entry (below).

